

lung-cancer-code

May 25, 2025

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
```

```
[2]: lung_data = pd.read_csv("c:/Users/shrut/Downloads/survey lung cancer.csv")
```

```
[3]: lung_data
```

```
[3]:
```

	GENDER	AGE	SMOKING	YELLOW_FINGERS	ANXIETY	PEER_PRESSURE	\
0	M	69	1	2	2	1	
1	M	74	2	1	1	1	
2	F	59	1	1	1	2	
3	M	63	2	2	2	1	
4	F	63	1	2	1	1	
..	
304	F	56	1	1	1	2	
305	M	70	2	1	1	1	
306	M	58	2	1	1	1	
307	M	67	2	1	2	1	
308	M	62	1	1	1	2	

	CHRONIC DISEASE	FATIGUE	ALLERGY	WHEEZING	ALCOHOL CONSUMING	COUGHING	\
0	1	2	1	2	2	2	
1	2	2	2	1	1	1	
2	1	2	1	2	1	2	
3	1	1	1	1	2	1	
4	1	1	1	2	1	2	
..	
304	2	2	1	1	2	2	
305	1	2	2	2	2	2	
306	1	1	2	2	2	2	
307	1	2	2	1	2	2	
308	1	2	2	2	2	1	

SHORTNESS OF BREATH	SWALLOWING DIFFICULTY	CHEST PAIN	LUNG_CANCER
---------------------	-----------------------	------------	-------------

0		2		2		2	YES
1		2		2		2	YES
2		2		1		2	NO
3		1		2		2	NO
4		2		1		1	NO
..		
304		2		2		1	YES
305		2		1		2	YES
306		1		1		2	YES
307		2		1		2	YES
308		1		2		1	YES

[309 rows x 16 columns]

```
[4]: lung_data.head()
```

```
[4]:  GENDER  AGE  SMOKING  YELLOW_FINGERS  ANXIETY  PEER_PRESSURE  \
0      M    69         1                2         2             1
1      M    74         2                1         1             1
2      F    59         1                1         1             2
3      M    63         2                2         2             1
4      F    63         1                2         1             1

      CHRONIC DISEASE  FATIGUE  ALLERGY  WHEEZING  ALCOHOL CONSUMING  COUGHING  \
0                    1        2        1         2                2         2
1                    2        2        2         1                1         1
2                    1        2        1         2                1         2
3                    1        1        1         1                2         1
4                    1        1        1         2                1         2

      SHORTNESS OF BREATH  SWALLOWING DIFFICULTY  CHEST PAIN  LUNG_CANCER
0                        2                      2          2         YES
1                        2                      2          2         YES
2                        2                      1          2         NO
3                        1                      2          2         NO
4                        2                      1          1         NO
```

```
[5]: lung_data.tail()
```

```
[5]:  GENDER  AGE  SMOKING  YELLOW_FINGERS  ANXIETY  PEER_PRESSURE  \
304    F    56         1                1         1             2
305    M    70         2                1         1             1
306    M    58         2                1         1             1
307    M    67         2                1         2             1
308    M    62         1                1         1             2

      CHRONIC DISEASE  FATIGUE  ALLERGY  WHEEZING  ALCOHOL CONSUMING  COUGHING  \
```

304	2	2	1	1	2	2
305	1	2	2	2	2	2
306	1	1	2	2	2	2
307	1	2	2	1	2	2
308	1	2	2	2	2	1

	SHORTNESS OF BREATH	SWALLOWING DIFFICULTY	CHEST PAIN	LUNG_CANCER
304	2	2	1	YES
305	2	1	2	YES
306	1	1	2	YES
307	2	1	2	YES
308	1	2	1	YES

```
[6]: #dependent_variable
x = lung_data.iloc[:,0:-1]
print(x)
```

	GENDER	AGE	SMOKING	YELLOW_FINGERS	ANXIETY	PEER_PRESSURE	\
0	M	69	1	2	2	1	
1	M	74	2	1	1	1	
2	F	59	1	1	1	2	
3	M	63	2	2	2	1	
4	F	63	1	2	1	1	
..	
304	F	56	1	1	1	2	
305	M	70	2	1	1	1	
306	M	58	2	1	1	1	
307	M	67	2	1	2	1	
308	M	62	1	1	1	2	

	CHRONIC DISEASE	FATIGUE	ALLERGY	WHEEZING	ALCOHOL CONSUMING	COUGHING	\
0	1	2	1	2	2	2	
1	2	2	2	1	1	1	
2	1	2	1	2	1	2	
3	1	1	1	1	2	1	
4	1	1	1	2	1	2	
..	
304	2	2	1	1	2	2	
305	1	2	2	2	2	2	
306	1	1	2	2	2	2	
307	1	2	2	1	2	2	
308	1	2	2	2	2	1	

	SHORTNESS OF BREATH	SWALLOWING DIFFICULTY	CHEST PAIN
0	2	2	2
1	2	2	2
2	2	1	2

3	1	2	2
4	2	1	1
..
304	2	2	1
305	2	1	2
306	1	1	2
307	2	1	2
308	1	2	1

[309 rows x 15 columns]

```
[7]: #independent_variable
y = lung_data. iloc[:,-1:]
print(y)
```

	LUNG_CANCER
0	YES
1	YES
2	NO
3	NO
4	NO
..	...
304	YES
305	YES
306	YES
307	YES
308	YES

[309 rows x 1 columns]

```
[8]: lung_data.GENDER = lung_data.GENDER.map({"M":1,"F":2})
lung_data.LUNG_CANCER = lung_data.LUNG_CANCER.map({"YES":1,"NO":2})
```

```
[9]: lung_data.shape
```

```
[9]: (309, 16)
```

```
[10]: lung_data.isnull().sum()
```

```
[10]: GENDER          0
AGE                0
SMOKING            0
YELLOW_FINGERS     0
ANXIETY            0
PEER_PRESSURE      0
CHRONIC_DISEASE    0
FATIGUE            0
```

```

ALLERGY          0
WHEEZING         0
ALCOHOL CONSUMING 0
COUGHING         0
SHORTNESS OF BREATH 0
SWALLOWING DIFFICULTY 0
CHEST PAIN       0
LUNG_CANCER      0
dtype: int64

```

```
[11]: lung_data.dtypes
```

```

[11]: GENDER          int64
      AGE            int64
      SMOKING         int64
      YELLOW_FINGERS  int64
      ANXIETY         int64
      PEER_PRESSURE   int64
      CHRONIC DISEASE int64
      FATIGUE         int64
      ALLERGY         int64
      WHEEZING        int64
      ALCOHOL CONSUMING int64
      COUGHING        int64
      SHORTNESS OF BREATH int64
      SWALLOWING DIFFICULTY int64
      CHEST PAIN      int64
      LUNG_CANCER     int64
      dtype: object

```

```

[12]: #the describe() method returns description of data in DataFrame
      lung_data.describe()

```

```

[12]:
count  GENDER  AGE  SMOKING  YELLOW_FINGERS  ANXIETY  \
count  309.000000  309.000000  309.000000  309.000000  309.000000
mean    1.475728  62.673139  1.563107  1.569579  1.498382
std     0.500221  8.210301  0.496806  0.495938  0.500808
min     1.000000  21.000000  1.000000  1.000000  1.000000
25%     1.000000  57.000000  1.000000  1.000000  1.000000
50%     1.000000  62.000000  2.000000  2.000000  1.000000
75%     2.000000  69.000000  2.000000  2.000000  2.000000
max     2.000000  87.000000  2.000000  2.000000  2.000000

      PEER_PRESSURE  CHRONIC DISEASE  FATIGUE  ALLERGY  WHEEZING  \
count  309.000000  309.000000  309.000000  309.000000  309.000000
mean    1.501618  1.504854  1.673139  1.556634  1.556634
std     0.500808  0.500787  0.469827  0.497588  0.497588

```

min	1.000000	1.000000	1.000000	1.000000	1.000000
25%	1.000000	1.000000	1.000000	1.000000	1.000000
50%	2.000000	2.000000	2.000000	2.000000	2.000000
75%	2.000000	2.000000	2.000000	2.000000	2.000000
max	2.000000	2.000000	2.000000	2.000000	2.000000

	ALCOHOL CONSUMING	COUGHING	SHORTNESS OF BREATH \
count	309.000000	309.000000	309.000000
mean	1.556634	1.579288	1.640777
std	0.497588	0.494474	0.480551
min	1.000000	1.000000	1.000000
25%	1.000000	1.000000	1.000000
50%	2.000000	2.000000	2.000000
75%	2.000000	2.000000	2.000000
max	2.000000	2.000000	2.000000

	SWALLOWING DIFFICULTY	CHEST PAIN	LUNG_CANCER
count	309.000000	309.000000	309.000000
mean	1.469256	1.556634	1.126214
std	0.499863	0.497588	0.332629
min	1.000000	1.000000	1.000000
25%	1.000000	1.000000	1.000000
50%	1.000000	2.000000	1.000000
75%	2.000000	2.000000	1.000000
max	2.000000	2.000000	2.000000

```
[13]: #the info() method prints information of the database
lung_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 309 entries, 0 to 308
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   GENDER                309 non-null    int64
1   AGE                   309 non-null    int64
2   SMOKING               309 non-null    int64
3   YELLOW_FINGERS       309 non-null    int64
4   ANXIETY               309 non-null    int64
5   PEER_PRESSURE        309 non-null    int64
6   CHRONIC_DISEASE      309 non-null    int64
7   FATIGUE              309 non-null    int64
8   ALLERGY              309 non-null    int64
9   WHEEZING             309 non-null    int64
10  ALCOHOL_CONSUMING     309 non-null    int64
11  COUGHING              309 non-null    int64
12  SHORTNESS OF BREATH  309 non-null    int64
```

```

13 SWALLOWING DIFFICULTY 309 non-null int64
14 CHEST PAIN             309 non-null int64
15 LUNG_CANCER            309 non-null int64
dtypes: int64(16)
memory usage: 38.8 KB

```

```

[14]: #Splitting the Dataset: Training and Testing
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=1/
↪3,random_state=0)

```

```

[15]: lung_data['LUNG_CANCER'].value_counts()

```

```

[15]: LUNG_CANCER
1      270
2       39
Name: count, dtype: int64

```

```

[16]: len(lung_data)

```

```

[16]: 309

```

```

[17]: #dependent_variable
x = lung_data.iloc[:,0:-1]
x

```

```

[17]:
   GENDER  AGE  SMOKING  YELLOW_FINGERS  ANXIETY  PEER_PRESSURE  \
0         1   69         1                2         2             1
1         1   74         2                1         1             1
2         2   59         1                1         1             2
3         1   63         2                2         2             1
4         2   63         1                2         1             1
..      ...  ...      ...              ...      ...             ...
304        2   56         1                1         1             2
305        1   70         2                1         1             1
306        1   58         2                1         1             1
307        1   67         2                1         2             1
308        1   62         1                1         1             2

   CHRONIC DISEASE  FATIGUE  ALLERGY  WHEEZING  ALCOHOL CONSUMING  COUGHING  \
0                 1         2         1         2                 2         2
1                 2         2         2         1                 1         1
2                 1         2         1         2                 1         2
3                 1         1         1         1                 2         1
4                 1         1         1         2                 1         2
..              ...      ...      ...      ...              ...      ...
304              2         2         1         1                 2         2

```

305	1	2	2	2	2	2
306	1	1	2	2	2	2
307	1	2	2	1	2	2
308	1	2	2	2	2	1

	SHORTNESS OF BREATH	SWALLOWING DIFFICULTY	CHEST PAIN
0	2	2	2
1	2	2	2
2	2	1	2
3	1	2	2
4	2	1	1
..
304	2	2	1
305	2	1	2
306	1	1	2
307	2	1	2
308	1	2	1

[309 rows x 15 columns]

```
[18]: #independent_variable
y = lung_data.iloc[:,-1:]
y
```

```
[18]: LUNG_CANCER
0      1
1      1
2      2
3      2
4      2
..
304    1
305    1
306    1
307    1
308    1
```

[309 rows x 1 columns]

```
[19]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
```



```
[20]: from sklearn.linear_model import LogisticRegression
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=1/
↳3,random_state=0)
```

```
[21]: #Fitting simple linear regression to the training test
Model1 = LogisticRegression()
Model1.fit(x_train, y_train)
#Predicting the test set results
prediction1 = Model1.predict(x_test)
```

C:\Users\shrut\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\sklearn\utils\validation.py:1408: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

C:\Users\shrut\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\sklearn\linear_model_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
[22]: prediction1
```

```
[22]: array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 2, 1, 1,
1, 1, 2, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1], dtype=int64)
```

```
[23]: from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
confusion_matrix(y_test,prediction1)
```

```
[23]: array([[85,  2],
[10,  6]], dtype=int64)
```

```
[24]: accuracy_score(y_test,prediction1)
```

```
[24]: 0.883495145631068
```

```
[25]: from sklearn.metrics import precision_score
      probs = Model1.predict_proba(x_test)
      precision_score(y_test, prediction1, average = None)
```

```
[25]: array([0.89473684, 0.75      ])
```

```
[26]: from sklearn.metrics import precision_score, recall_score, f1_score

      # assuming your predicted and actual labels are stored in variables y_pred and
      # y_true, respectively
      accuracy = accuracy_score(y_test, prediction1)
      precision = precision_score(y_test, prediction1)
      recall = recall_score(y_test, prediction1)
      f1 = f1_score(y_test, prediction1)

      print("Accuracy:", accuracy)
      print("Precision:", precision)
      print("Recall:", recall)
      print("F1 score:", f1)
```

```
Accuracy: 0.883495145631068
Precision: 0.8947368421052632
Recall: 0.9770114942528736
F1 score: 0.9340659340659341
```

```
[27]: from sklearn.metrics import recall_score
      from sklearn.metrics import f1_score
```


```
[28]: recall_score(y_test, prediction1, average = None)
```

```
[28]: array([0.97701149, 0.375      ])
```

```
[29]: f1_score(y_test, prediction1, average = None)
```

```
[29]: array([0.93406593, 0.5       ])
```

```
[30]: cm = confusion_matrix(y_true = y_test, y_pred = prediction1)
      #plot_confusion_matrix(cm, level, title = "confusion_matrix")
      sns.heatmap(cm, annot=True, cmap="Blues", fmt="d")
      plt.xlabel("Predicted Labels")
      plt.ylabel("True Labels")
      plt.title("Confusion Matrix")
      plt.show()
```



```
[31]: from sklearn.neighbors import KNeighborsClassifier
```

```
[32]: #Fitting K-NN to the Training set
classifier = KNeighborsClassifier(n_neighbors = 3, metric = "minkowski", p = 2)
classifier.fit(x_train, y_train)
```

C:\Users\shrut\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\sklearn\neighbors_classification.py:239: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return self._fit(X, y)
```

```
[32]: KNeighborsClassifier(n_neighbors=3)
```

```
[33]: #Predicting the Test set result
prediction2 = classifier.predict(x_test)
```

```
[34]: prediction2
```

```
[34]: array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1,
          1, 1, 2, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
          1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
          1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1,
          1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1], dtype=int64)
```

```
[35]: from sklearn.metrics import confusion_matrix
      from sklearn.metrics import accuracy_score
      confusion_matrix(y_test, prediction2)
```

```
[35]: array([[86,  1],
          [12,  4]], dtype=int64)
```

```
[36]: from sklearn.metrics import precision_score, recall_score, f1_score

      # assuming your predicted and actual labels are stored in variables y_pred and
      # y_true, respectively
      accuracy = accuracy_score(y_test, prediction2)
      precision = precision_score(y_test, prediction2)
      recall = recall_score(y_test, prediction2)
      f1 = f1_score(y_test, prediction2)

      print("Accuracy:", accuracy)
      print("Precision:", precision)
      print("Recall:", recall)
      print("F1 score:", f1)
```

```
Accuracy: 0.8737864077669902
Precision: 0.8775510204081632
Recall: 0.9885057471264368
F1 score: 0.9297297297297298
```

```
[37]: accuracy_score(y_test, prediction2)
```

```
[37]: 0.8737864077669902
```

```
[38]: probs = Model1.predict_proba(x_test)
      precision_score(y_test, prediction2, average = None)
```

```
[38]: array([0.87755102, 0.8       ])
```


```
[39]: recall_score(y_test, prediction2, average = None)
```

```
[39]: array([0.98850575, 0.25       ])
```

```
[40]: f1_score(y_test, prediction2, average = None)
```

```
[40]: array([0.92972973, 0.38095238])
```

```
[41]: cm = confusion_matrix(y_true = y_test, y_pred = prediction2)
      #plot_confusion_matrix(cm,level,title = "confusion_matrix")
      sns.heatmap(cm, annot=True, cmap="Blues", fmt="d")
      plt.xlabel("Predicted Labels")
      plt.ylabel("True Labels")
      plt.title("Confusion Matrix")
      plt.show()
```



```
[42]: #Decision Tree
from sklearn.tree import DecisionTreeClassifier
tree = DecisionTreeClassifier(random_state = 0,criterion = "entropy")
tree.fit(x_train, y_train)
prediction3 = classifier.predict(x_test)
```

```
[43]: prediction3
```

```
[43]: array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1,  
            1, 1, 2, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
```

```
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1], dtype=int64)
```

```
[44]: confusion_matrix(y_test, prediction3)
```

```
[44]: array([[86,  1],
          [12,  4]], dtype=int64)
```

```
[45]: from sklearn.metrics import precision_score, recall_score, f1_score

# assuming your predicted and actual labels are stored in variables y_pred and
# y_true, respectively
accuracy = accuracy_score(y_test, prediction3)
precision = precision_score(y_test, prediction3)
recall = recall_score(y_test, prediction3)
f1 = f1_score(y_test, prediction3)

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 score:", f1)
```

```
Accuracy: 0.8737864077669902
Precision: 0.8775510204081632
Recall: 0.9885057471264368
F1 score: 0.9297297297297298
```

```
[46]: accuracy_score(y_test, prediction3)
```

```
[46]: 0.8737864077669902
```

```
[47]: probs = Model1.predict_proba(x_test)
precision_score(y_test, prediction3, average = None)
```

```
[47]: array([0.87755102, 0.8       ])
```

```
[48]: recall_score(y_test, prediction3, average = None)
```

```
[48]: array([0.98850575, 0.25       ])
```

```
[49]: f1_score(y_test, prediction3, average = None)
```


```
[49]: array([0.92972973, 0.38095238])
```

```
[50]: cm = confusion_matrix(y_true = y_test, y_pred = prediction3)
#plot_confusion_matrix(cm, level, title = "confusion_matrix")
```

```

sns.heatmap(cm, annot=True, cmap="Blues", fmt="d")
plt.xlabel("Predicted Labels")
plt.ylabel("True Labels")
plt.title("Confusion Matrix")
plt.show()

```



```

[51]: #Support Vector Machine
from sklearn.ensemble import BaggingClassifier
from sklearn.multiclass import OneVsRestClassifier
from sklearn.svm import SVC
svm =
    ↳OneVsRestClassifier(BaggingClassifier(SVC(C=10,kernel='rbf',random_state=9,probability=True
svm.fit(x_train, y_train)
prediction4 = svm.predict(x_test)

```

```

[52]: prediction4

```

```

[52]: array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
        1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
        1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
        1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,

```

```
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1], dtype=int64)
```

```
[53]: confusion_matrix(y_test, prediction4)
```

```
[53]: array([[87, 0],
             [16, 0]], dtype=int64)
```

```
[54]: from sklearn.metrics import precision_score, recall_score, f1_score

# assuming your predicted and actual labels are stored in variables y_pred and
# y_true, respectively
accuracy = accuracy_score(y_test, prediction4)
precision = precision_score(y_test, prediction4)
recall = recall_score(y_test, prediction4)
f1 = f1_score(y_test, prediction4)

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 score:", f1)
```

Accuracy: 0.8446601941747572
Precision: 0.8446601941747572
Recall: 1.0
F1 score: 0.9157894736842105

```
[55]: accuracy_score(y_test, prediction4)
```

```
[55]: 0.8446601941747572
```

```
[56]: probs = Model1.predict_proba(x_test)
precision_score(y_test, prediction4, average = None)
```

```
C:\Users\shrut\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
[56]: array([0.84466019, 0.          ])
```

```
[57]: recall_score(y_test, prediction4, average = None)
```


```
[57]: array([1., 0.])
```



```
[58]: f1_score(y_test, prediction4, average = None)
```

```
[58]: array([0.91578947, 0.          ])
```

```
[59]: cm = confusion_matrix(y_true = y_test, y_pred = prediction4)
#plot_confusion_matrix(cm, level, title = "confusion_matrix")
sns.heatmap(cm, annot=True, cmap="Blues", fmt="d")
plt.xlabel("Predicted Labels")
plt.ylabel("True Labels")
plt.title("Confusion Matrix")
plt.show()
```



```
[60]: from sklearn.naive_bayes import GaussianNB
nbcla = GaussianNB()
nbcla.fit(x_train, y_train)
prediction5 = nbcla.predict(x_test)
```

C:\Users\shrut\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\sklearn\utils\validation.py:1408: DataConversionWarning: A column-

vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

```
[61]: from sklearn.metrics import confusion_matrix
      from sklearn.metrics import accuracy_score
      confusion_matrix(y_test, prediction5)
```

```
[61]: array([[81,  6],
           [ 8,  8]], dtype=int64)
```

```
[62]: from sklearn.metrics import precision_score, recall_score, f1_score

      # assuming your predicted and actual labels are stored in variables y_pred and
      # y_true, respectively
      accuracy = accuracy_score(y_test, prediction5)
      precision = precision_score(y_test, prediction5)
      recall = recall_score(y_test, prediction5)
      f1 = f1_score(y_test, prediction5)

      print("Accuracy:", accuracy)
      print("Precision:", precision)
      print("Recall:", recall)
      print("F1 score:", f1)
```

Accuracy: 0.8640776699029126

Precision: 0.9101123595505618

Recall: 0.9310344827586207

F1 score: 0.9204545454545454

```
[63]: accuracy_score(y_test, prediction5)
```

```
[63]: 0.8640776699029126
```

```
[64]: probs = Model1.predict_proba(x_test)
      precision_score(y_test, prediction5, average = None)
```

```
[64]: array([0.91011236, 0.57142857])
```


```
[65]: recall_score(y_test, prediction5, average = None)
```

```
[65]: array([0.93103448, 0.5       ])
```

```
[66]: f1_score(y_test, prediction5, average = None)
```

```
[66]: array([0.92045455, 0.53333333])
```

```
[67]: cm = confusion_matrix(y_true = y_test, y_pred = prediction5)
#plot_confusion_matrix(cm, level, title = "confusion_matrix")
sns.heatmap(cm, annot=True, cmap="Blues", fmt="d")
plt.xlabel("Predicted Labels")
plt.ylabel("True Labels")
plt.title("Confusion Matrix")
plt.show()
```



```
[68]: from sklearn.ensemble import RandomForestClassifier

# Initialize the classifier
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)

# Train the model using training dataset
rf_classifier.fit(x_train, y_train)

# Make predictions on test dataset
prediction6 = rf_classifier.predict(x_test)

# Evaluate the accuracy of the model
```

```
#accuracy = rf_classifier.score(x_test, y_test)
#print("Accuracy:", accuracy)
```

C:\Users\shrut\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\sklearn\base.py:1389: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return fit_method(estimator, *args, **kwargs)
```

```
[69]: confusion_matrix(y_test, prediction6)
```

```
[69]: array([[85,  2],
          [ 9,  7]], dtype=int64)
```

```
[70]: from sklearn.metrics import precision_score, recall_score, f1_score
```

```
# assuming your predicted and actual labels are stored in variables y_pred and
↪ y_true, respectively
accuracy = accuracy_score(y_test, prediction6)
precision = precision_score(y_test, prediction6)
recall = recall_score(y_test, prediction6)
f1 = f1_score(y_test, prediction6)

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 score:", f1)
```

Accuracy: 0.8932038834951457
Precision: 0.9042553191489362
Recall: 0.9770114942528736
F1 score: 0.9392265193370166

```
[71]: accuracy_score(y_test, prediction6)
```

```
[71]: 0.8932038834951457
```

```
[72]: probs = Model1.predict_proba(x_test)
precision_score(y_test, prediction6, average = None)
```

```
[72]: array([0.90425532, 0.77777778])
```


```
[73]: recall_score(y_test, prediction6, average = None)
```

```
[73]: array([0.97701149, 0.4375    ])
```

```
[74]: f1_score(y_test, prediction6, average = None)
```

```
[74]: array([0.93922652, 0.56      ])
```

```
[75]: cm = confusion_matrix(y_true = y_test, y_pred = prediction6)
#plot_confusion_matrix(cm, level, title = "confusion_matrix")
sns.heatmap(cm, annot=True, cmap="Blues", fmt="d")
plt.xlabel("Predicted Labels")
plt.ylabel("True Labels")
plt.title("Confusion Matrix")
plt.show()
```



```
[76]: #Finding Correlation
cn=lung_data.corr()
cn
```

```
[76]:
```

	GENDER	AGE	SMOKING	YELLOW_FINGERS	ANXIETY \
GENDER	1.000000	-0.021306	-0.036277	0.212959	0.152127
AGE	-0.021306	1.000000	-0.084475	0.005205	0.053170
SMOKING	-0.036277	-0.084475	1.000000	-0.014585	0.160267

YELLOW_FINGERS	0.212959	0.005205	-0.014585	1.000000	0.565829
ANXIETY	0.152127	0.053170	0.160267	0.565829	1.000000
PEER_PRESSURE	0.275564	0.018685	-0.042822	0.323083	0.216841
CHRONIC DISEASE	0.204606	-0.012642	-0.141522	0.041122	-0.009678
FATIGUE	0.083560	0.012614	-0.029575	-0.118058	-0.188538
ALLERGY	-0.154251	0.027990	0.001913	-0.144300	-0.165750
WHEEZING	-0.141207	0.055011	-0.129426	-0.078515	-0.191807
ALCOHOL CONSUMING	-0.454268	0.058985	-0.050623	-0.289025	-0.165750
COUGHING	-0.133303	0.169950	-0.129471	-0.012640	-0.225644
SHORTNESS OF BREATH	0.064911	-0.017513	0.061264	-0.105944	-0.144077
SWALLOWING DIFFICULTY	0.078161	-0.001270	0.030718	0.345904	0.489403
CHEST PAIN	-0.362958	-0.018104	0.120117	-0.104829	-0.113634
LUNG_CANCER	0.067254	-0.089465	-0.058179	-0.181339	-0.144947

	PEER_PRESSURE	CHRONIC DISEASE	FATIGUE	ALLERGY \
GENDER	0.275564	0.204606	0.083560	-0.154251
AGE	0.018685	-0.012642	0.012614	0.027990
SMOKING	-0.042822	-0.141522	-0.029575	0.001913
YELLOW_FINGERS	0.323083	0.041122	-0.118058	-0.144300
ANXIETY	0.216841	-0.009678	-0.188538	-0.165750
PEER_PRESSURE	1.000000	0.048515	0.078148	-0.081800
CHRONIC DISEASE	0.048515	1.000000	-0.110529	0.106386
FATIGUE	0.078148	-0.110529	1.000000	0.003056
ALLERGY	-0.081800	0.106386	0.003056	1.000000
WHEEZING	-0.068771	-0.049967	0.141937	0.173867
ALCOHOL CONSUMING	-0.159973	0.002150	-0.191377	0.344339
COUGHING	-0.089019	-0.175287	0.146856	0.189524
SHORTNESS OF BREATH	-0.220175	-0.026459	0.441745	-0.030056
SWALLOWING DIFFICULTY	0.366590	0.075176	-0.132790	-0.061508
CHEST PAIN	-0.094828	-0.036938	-0.010832	0.239433
LUNG_CANCER	-0.186388	-0.110891	-0.150673	-0.327766


	WHEEZING	ALCOHOL CONSUMING	COUGHING \
GENDER	-0.141207	-0.454268	-0.133303
AGE	0.055011	0.058985	0.169950
SMOKING	-0.129426	-0.050623	-0.129471
YELLOW_FINGERS	-0.078515	-0.289025	-0.012640
ANXIETY	-0.191807	-0.165750	-0.225644
PEER_PRESSURE	-0.068771	-0.159973	-0.089019
CHRONIC DISEASE	-0.049967	0.002150	-0.175287
FATIGUE	0.141937	-0.191377	0.146856
ALLERGY	0.173867	0.344339	0.189524
WHEEZING	1.000000	0.265659	0.374265
ALCOHOL CONSUMING	0.265659	1.000000	0.202720
COUGHING	0.374265	0.202720	1.000000
SHORTNESS OF BREATH	0.037834	-0.179416	0.277385
SWALLOWING DIFFICULTY	0.069027	-0.009294	-0.157586

CHEST PAIN	0.147640	0.331226	0.083958
LUNG_CANCER	-0.249300	-0.288533	-0.248570

	SHORTNESS OF BREATH	SWALLOWING DIFFICULTY	CHEST PAIN \
GENDER	0.064911	0.078161	-0.362958
AGE	-0.017513	-0.001270	-0.018104
SMOKING	0.061264	0.030718	0.120117
YELLOW_FINGERS	-0.105944	0.345904	-0.104829
ANXIETY	-0.144077	0.489403	-0.113634
PEER_PRESSURE	-0.220175	0.366590	-0.094828
CHRONIC DISEASE	-0.026459	0.075176	-0.036938
FATIGUE	0.441745	-0.132790	-0.010832
ALLERGY	-0.030056	-0.061508	0.239433
WHEEZING	0.037834	0.069027	0.147640
ALCOHOL CONSUMING	-0.179416	-0.009294	0.331226
COUGHING	0.277385	-0.157586	0.083958
SHORTNESS OF BREATH	1.000000	-0.161015	0.024256
SWALLOWING DIFFICULTY	-0.161015	1.000000	0.069027
CHEST PAIN	0.024256	0.069027	1.000000
LUNG_CANCER	-0.060738	-0.259730	-0.190451

	LUNG_CANCER
GENDER	0.067254
AGE	-0.089465
SMOKING	-0.058179
YELLOW_FINGERS	-0.181339
ANXIETY	-0.144947
PEER_PRESSURE	-0.186388
CHRONIC DISEASE	-0.110891
FATIGUE	-0.150673
ALLERGY	-0.327766
WHEEZING	-0.249300
ALCOHOL CONSUMING	-0.288533
COUGHING	-0.248570
SHORTNESS OF BREATH	-0.060738
SWALLOWING DIFFICULTY	-0.259730
CHEST PAIN	-0.190451
LUNG_CANCER	1.000000

```
[77]: #Correlation
cmap=sns.diverging_palette(260,-10,s=50, l=75, n=6,
as_cmap=True)
plt.subplots(figsize=(18,18))
sns.heatmap(cn,cmap="Blues",annot=True, square=True)
plt.show()
```

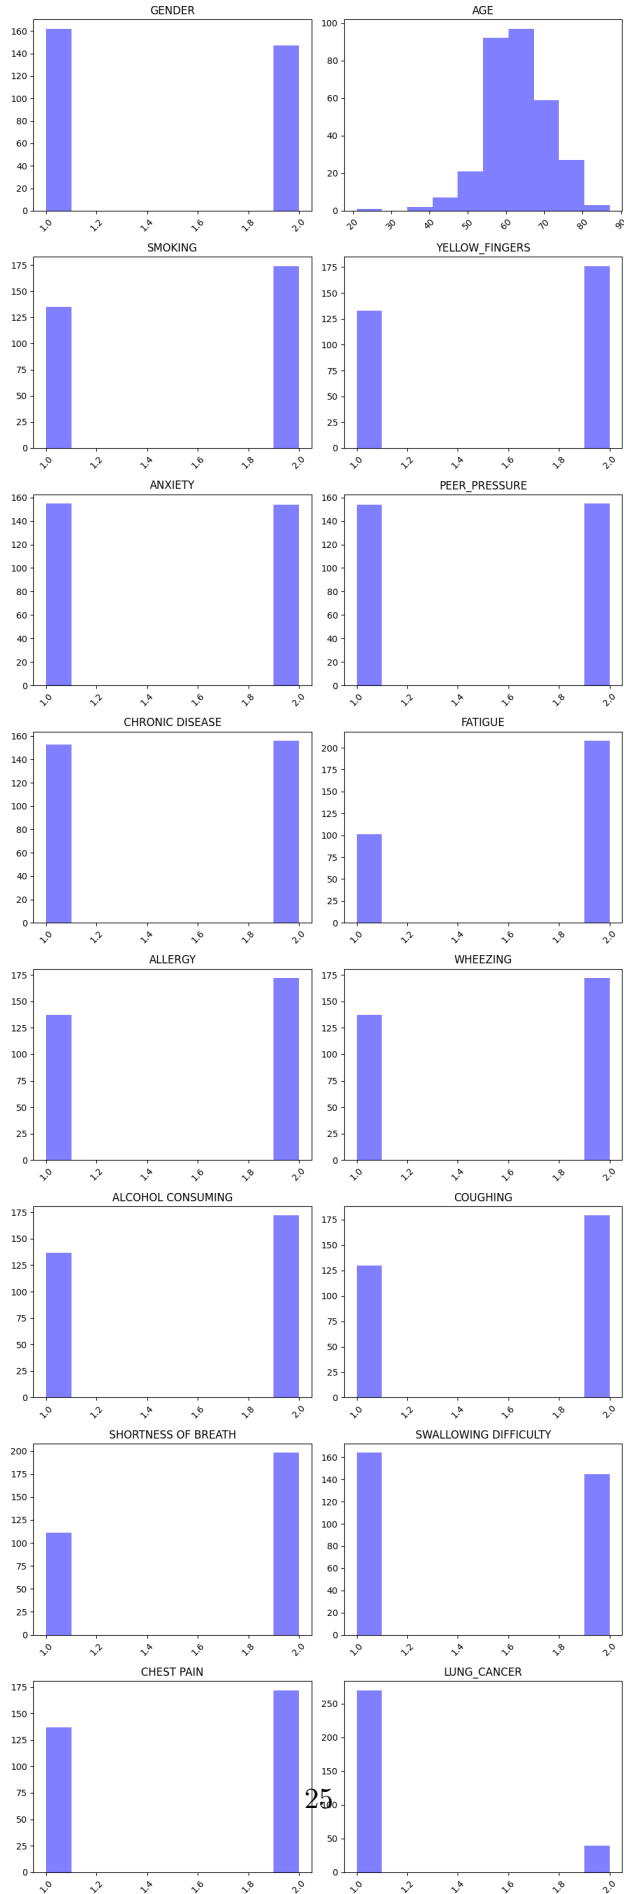


```
[78]: num_list = list(lung_data.columns)

fig = plt.figure(figsize=(10,30))

for i in range(len(num_list)):
    plt.subplot(8,2,i+1)
    plt.title(num_list[i])
    plt.xticks(rotation=45)
    plt.hist(lung_data[num_list[i]],color='blue',alpha=0.5)

plt.tight_layout()
```

[]: